

Youngjoon Cheon

Department of Electrical and Computer Engineering, Korea University

yjcheon@csl.korea.ac.kr ◇ cheonyj7@gmail.com

Homepage ◇ Google Scholar

RESEARCH INTERESTS

- **Accelerator Memory Systems & Heterogeneous Architecture:** Architectures for expanding usable memory capacity and bandwidth in data-intensive workloads
- **Hardware–Software Co-Designed Memory Management:** Workload-aware mechanisms for data placement, movement, and address translation

EDUCATION

Korea University

M.S. in Electrical and Computer Engineering, GPA: 3.94/4.5

Thesis: *Segmentation-based Address Translation in GPUs for Accelerating LLM Inference*

Advisor: Prof. Jeongseob Ahn

Sep. 2024 – Aug. 2026

Seoul, South Korea

Ajou University

B.S. in Software and Computer Engineering, GPA: 3.77/4.5

Minor in Mathematics

Mar. 2018 – Aug. 2024

Suwon, South Korea

RESEARCH EXPERIENCE

Graduate Research Assistant

Computer Systems Laboratory, Korea University

Sep. 2024 – Aug. 2026

Advisor: Prof. Jeongseob Ahn

- **HW/SW Co-Design for LLM Address Translation [CAL’2026]** Jan. 2025 – May. 2026
 - Identified KV-cache-induced TLB contention that evicted reused model-weight translations during long-context LLM inference
 - Co-designed vLLM weight placement and a GPU segmentation table to bypass TLB translation for static, read-only weights
 - Reduced the L2 TLB miss rate from 91.5% to 17.0% and achieved 2.51× the IPC of the paging baseline in Accel-Sim
- **Intra-GPU Memory Tiering for KV-Cache Scaling** Apr. 2024 – Dec. 2024
 - Designed an HBM–LPDDR hierarchy with separate controllers, a unified address space, and layer-granular KV-cache migration
 - Integrated Ramulator with GPGPU-Sim to model distinct HBM/LPDDR timing behavior
 - Identified simulator-induced SM contention during migration and added a DMA copy-engine model to decouple data movement from computation
- **Undergraduate Research Intern** Sep. 2022 – Aug. 2024

Computer Systems Laboratory, Ajou University Advisor: Prof. Jeongseob Ahn
- **Heterogeneous xPU–PIM Execution for LLM Inference** Jul. 2023 – Mar. 2024
 - Partitioned Hugging Face autoregressive decoding to offload memory-bound attention over the KV-cache to emulated PIM while retaining QKV projections and feed-forward layers on the xPU
 - Built a multi-GPU PIM emulation testbed and profiled inter-device overheads with NVIDIA Nsight Systems
 - Doubled memory capacity and supported batch size at 3.4% higher end-to-end latency than the single-GPU baseline

· **System Software & Parallel Programming Studies**

Sep. 2022 – Jun. 2023

- Analyzed Data Access MONitor (DAMON) subsystem internals in the Linux kernel
- Investigated CUDA and parallel programming models for GPU accelerator research

PUBLICATIONS

Youngjoon Cheon, Yunho Oh, and Jeongseob Ahn, “Towards Segmentation-based Address Translation for LLM Inference,” *IEEE Computer Architecture Letters (CAL)*, 2026. [Paper]

AWARDS AND HONORS

Academic Excellence Scholarship (30-50% Tuition Waivers) Fall 2019, Fall 2022, Fall 2023
Ajou University

TEACHING EXPERIENCE

Graduate Teaching Assistant, Korea University

- **Digital Systems and Design Laboratory** Fall 2024, Fall 2025
 - Conducted weekly lectures on digital logic and guided hands-on Verilog/FPGA synthesis labs
- **Computer Architecture** Spring 2025, Spring 2026
 - Developed a RISC-V assembler programming assignment ([GitHub link](#))

TECHNICAL SKILLS

Languages	Verilog, C/C++, CUDA, Python, Assembly, Shell
Frameworks	PyTorch, vLLM
Simulators & Tools	GPGPU-Sim, Accel-Sim, Ramulator, DRAMSim2, Vivado, Git

PROFESSIONAL SERVICE

Artifact Evaluation Committee

- ACM Symposium on Operating Systems Principles (SOSP 2026)
- IEEE International Symposium on High-Performance Computer Architecture (HPCA 2026)